

Programming in Lua Notes

Programming in Lua gives a solid base for any programmer who wants to use Lua. It covers all aspects of Lua—from the basics to its API with C—explaining how to make good use of its features and giving numerous code examples. The book is targeted at people with some programming background, but it does not assume any prior knowledge about Lua or other scripting languages.

Getting Started

- Each piece of code that Lua executes is called a *chunk*
- Lua needs no separator between consecutive statements (but you can use a semicolon if you wish)
- To exit the interpreter, use `os.exit()`
- To parse a file from Lua, use `dofile("filename.lua")`
- Identifiers can be any string of letters, digits and underscores not beginning with a digit
- Lua 5.2 accepts only English letters for identifiers (`a-z` and `A-Z`)
- Comments start with a double hyphen (`--`) and go to the rest of the line
- Long comments start with a `--[[` and end with `]]`
- Global variables do not need declarations, they are `nil` by default
- In interactive mode, prepending an equals sign (`=`) to any expression prints the result of that expression
- Any arguments to a script are in the global variable `arg` by default

Types and Values

- Lua has eight basic types: `nil`, `boolean`, `number`, `string`, `userdata`, `function`, `thread` and `table`
- The type of a variable can be checked using the `type()` function, which returns a string representing the type of the given variable
- Functions are *first-class values* in Lua, they can be used like any other type of variables
- Lua uses `nil` as a kind of non-value, representing the absence of a useful value
- **All** numbers in Lua are real (double) *floating-point numbers* (there is no integer type)
- In the Lua number type, any integer up to 2^{53} has an exact representation
- Due to using a double type, there can be **rounding errors**: `12.7-20+7.3` is not exactly zero because both 12.7 and 7.3 do not have an *exact representation*
- Number literals can be written with both an optional floating-point part (eg. `0.1212`) and exponent (eg. `3.4e-4`), and can be either in base ten or

hexadecimal (with the 0x prefix, eg. 0xFFFF)

- Strings can contain any characters (null, any UTF-8 characters, etc.)
- Strings are *immutable* in Lua (cannot be modified)
- The length of a string can be acquired with the *length operator* (#)
- Strings can be delimited by single and double quotation marks ('str' and "str") as well as with double square brackets ([[str]])
- Strings can be concatenated with double periods (..)
- Strings and numbers can be converted with `tostring()` and `tonumber()`
- The only real data type in Lua are tables, which can be used to construct *arrays* (sequences) as well as *records*
- Tables are handled **by reference** (thus, `{ } ~= { }`)
- To access the member `abc` of table `t`, both `t["abc"]` and `t.abc` can be used
- Lua global variables are stored in a table
- Lua arrays are tables that use numbers from 1 to `n` as indexes, where `n` is the length of the array
- The length of Lua arrays without *holes* (embedded nils) can be acquired using the *length operator* (#)
- Userdata variables allow C data to be stored in Lua variables

Expressions

- Exponentiation is done in Lua with the caret (^)
- Modulus is obtained from a number with the percent sign (%)
- The fractional and integer part of a number can be obtained using the modulus operator (`n%1` for the former and `n-n%1` for the latter)
- The negation of the equality operator in Lua is a tilde and an equals sign combined (`~=`)
- Tables and userdata are compared **by reference**
- Strings are compared in **alphabetical order** (as determined by the locale)
- Logical operators (**not**, **and**, **or**) use *short-cut evaluation*, so `f()` or `error()` is only going to call `error()` if `f()` returns false
- The Lua idiom `x = x or v` sets `x` to `v` only if `x` is not `nil` or `false`
- The Lua idiom `c and t or f` returns `t` if `c` evaluates to `true`, and `f` otherwise (unless `t` evaluates to `false`)
- Concatenation in Lua is done with two dots (..). If one of the operands is a number, it is converted to a string automatically
- The *concatenation operator* does not modify its operands

- The length operator (#) works on strings and tables, on the latter it gives the length of the sequence represented by it, a sequence being a table where numeric keys go from 1 to n without any holes (embedded nils)
- Tables can be constructed by a few different constructors:
 - List constructor** constructs the table to be a sequence, has no notion of keys and looks like this: {324, "value two", true, ...}
 - Record constructor** constructs the table to be a record, has keys which must be valid Lua identifiers and looks like this: {fieldone=10, fieldtwo="value two", fieldthree=false, ...}
 - General constructor** can construct any kind of table, it's keys do not need to be valid identifiers and can in fact be of any type, it looks like this: {["field one"]=324, ["field two"]="value two", ...}

Statements

- Lua allows *multiple assignment*, which assigns a list of values to a list of variables in one step, both lists have their elements separated by commas
- Lua first evaluates all values and only then executes the assignments (allowing us to swap two variables with multiple assignment)
- When there are more variables than values, they are filled with nils
- When there are more values than variables, they are silently discarded
- A frequent use of multiple assignment is to collect *multiple returns* from function calls
- Lua supports local variables with the keyword `local`
- In interactive mode, **local variables don't work as expected** because every line is executed in its own chunk
- To make local variables work in interactive mode, they need to be put into a *do-end block*
- Access to local variables is faster than to global ones
- A common idiom in Lua is `local foo = foo`, which creates a local variable `foo` and assigns the global variable `foo` to it
- Lua supports *if-then-[/elseif-then]-else-end*, *while-do*, *repeat-until*, *numeric for* and *generic for* control structures
- *Numeric for* starts at a given start value and ends at a given end value using the steps provided: `for <var> = <start>, <end>, [<step=1>] do...`
- The value of the control variable should not be changed (use `break` to prematurely exit the loop)
- The *generic for* traverses all values returned by an iterator function, like so: `for key, value in pairs(table) do...`
- There are several *iterators*: `pairs()` to traverse a table, `io.lines()` to iterate over the lines of a file, `ipairs()` to iterate over the entries of a sequence, `string.gmatch()` to iterate over words in a string and more

- A return statement can only appear as the last statement of a block
- Lua supports **goto** and labels, they are declared with `::labelname::` and can be jumped to with `goto labelname`
- You cannot jump into a block, out of a function or into the scope of a local variable
- The scope of a local variable ends on the last *non-void statement* of the block where the variable is defined, labels are considered void statements
- Gotos can be used to emulate functionality like **continue**

Functions

- If a function has one single argument and that argument is either a string literal or a table constructor, the parentheses (in a function call) are optional
- The colon operator in Lua offers special syntax for *object oriented programming*: `o:method(a,b)` translates to `o.method(o,a,b)`
- You can call a function with a number of arguments different from it's number of parameters: extra arguments are thrown away, missing ones filled with **nil**
- Functions in Lua can return multiple results
- In some cases, like when the function is placed in parentheses (like so: `(f())`) or when it's used as an expression, only the first result is used
- The Lua function `table.unpack()` takes an array as input and returns the contents (using multiple return values)
- The opposite can be done with the function `table.pack()`, which turns all of it's parameter into an array and additionally stores the size of that array in the field **n**
- Lua functions can take a variable amount of inputs with the *vararg expression* `(...)`, which is used in place of the parameter list and expands to the given arguments in the function body
- *Named parameters* can be simulated in Lua by passing a table as the first and only argument, which allows us to have on key/value pair per argument, where the key is the name of that argument
- Can look like this: `copy{src="file1", dest="file2"}`

More about Functions

- Functions in Lua are *first-class values* with *proper lexical scoping*, meaning that they can access variables of their enclosing functions
- Functions can be stored in tables, and passed to and returned from other functions
- Functions are *anonymous* (not bound to any name)
- A Function definition is actually an assignment

- Functions as first-class values can be used to write *callback functions* or provide a sorting strategy to `table.sort()`
- Functions that get other functions as an argument are called *higher-order functions*
- The variables of the parent function that a function defined inside it can access are neither local nor global variables, these are called *nonlocal variables* or *upvalues* (they *escape* their original scope)
- *Closures* make use of proper lexical scoping: they are functions with access to nonlocal variables
- *Nonlocal variables* persist between function calls, similar to `static` function variables in C
- Closures can be used to create *sandboxes* by redefining functions in a more limited manner and hiding the original functions
- Functions can also be stored in local variables, and Lua has syntactic sugar to do this (by prepending `local` before a function declaration)
- When using indirect local recursive functions, they need a kind of *forward declaration* to indicate that they will be local (with `local name`) and they then need to be defined without the local function syntactic sugar
- Lua does proper *tail-call elimination* (tail calls do not cost stack space)
- Tail calls need to be in the form `return func(args)`

Iterators and the Generic for

- An iterator is any construction that allows you to *iterate* over the elements of a collection
- They are typically represented by functions (closures) in Lua
- A closure iterator involves two functions: the closure itself and a *factory*, which creates the closure and its nonlocal variables (the *state*)
- Iterators may not be easy to write, but they are easy to use
- The generic `for` does all the bookkeeping for an iteration loop and it also keeps an *invariant state variable* (can be used to keep a state) and a *control variable*
- When the first variable returned by the iterator (called the control variable) is `nil`, the loop ends
- With the invariant state and the control variable, we can write *stateless iterators* (like `ipairs()`, which is also stateless): these do not use nonlocal variables to keep their state
- Complex states can be stored in the invariant state variable by using a table
- *True iterators* are functions that do the iteration themselves, they take an anonymous function as argument and call that for every element
- They aren't used very much anymore since they have some drawbacks (like difficult parallel iteration)

Compilation, Execution and Errors

- Lua always *precompiles* source code to intermediate form (bytecode) before running it
- Lua is still considered an *interpreted language* since it is possible to execute code generated on the fly (with functions such as `load()`)
- The function `loadfile()` loads a Lua chunk from a file and returns a function that will call the chunk if called, or an error code
- We can use `loadfile()` to run a file several times (by executing the returned function multiple times)
- The `load()` function is similar, but it reads its chunk from a string
- The `load()` function is powerful and rather expensive, so it should be used **with care** and only when needed
- `load()` compiles code in the global environment, **without lexical scoping**
- You can use vararg expressions in `load()`ed strings since the code is treated as an anonymous function
- The `string.rep()` function repeats a string a given number of times
- `load()` can take a reader function as argument, which returns the chunk in parts
- `io.lines(filename, "*L")` returns a function that iterates over the lines in the given file
- `io.lines(filename, 1024)` is more efficient since it uses a fixed-size buffer
- The `load()` and `loadfile()` functions **never** have any side effects
- External chunks should be run in a *protected environment*
- Lua allows code to be distributed in precompiled form, such code is allowed anywhere normal code would be allowed as well
- Code can be precompiled with the `luac` program
- `string.dump()` returns the precompiled code (as a string) of any Lua function
- **Maliciously corrupted binary code can crash the Lua interpreter or even execute user-provided machine code!**
- As a second parameter, `load()` can accept a name of the chunk to be loaded for debugging purposes
- The third parameter to `load()` controls what kind of chunks can be loaded ('t' for textual, 'b' for binary and 'bt' for both)
- Lua supports *dynamic linking* even though that is not standard ANSI C
- To dynamically link to a library, use `package.loadlib(libpath, funcname)`, which returns the requested function
- Often libraries are loaded with `require()`, which auto-imports all functions and puts them into a package
- Whenever an error is raised, Lua ends the current chunk and returns to the application
- The `assert()` functions checks if its first argument is not false, if so it

returns it, else it raises an error

- Functions can return `false` and an error code to show errors or call the `error()` function directly
- Most functions return `false` and an error code so the error can be handled
- Errors raised with `error()` can be caught using the `pcall()` function, which stands for *protected call*
- `pcall()` takes a function to be called in protected mode as well as a level argument to tell which of the functions in the call stack is the culprit
- If we want a traceback of the error, we can use the `xpcall()` function, which takes a *message handler function* (which is called before the stack unwinds)
- Two common message handlers are `debug.debug` (provides interactive console) and `debug.traceback` (builds an extended error message with the traceback)

Coroutines

- Coroutines in Lua are like threads: they are a line of execution with their own stack, local variables and instruction pointer but sharing the global variables
- Coroutines run *concurrently*, not *parallel*: there's always **just one** coroutine currently running
- Coroutines are a means of *cooperative multitasking* (as opposed to *preemptive multitasking*): their execution is only suspended if they explicitly ask for it
- All coroutine functions are in the `coroutine` table
- They can be created with `coroutine.create()`, which takes a function as argument
- Coroutines are of type `thread`
- Coroutines can be in one of four states:

normal This is the state a coroutine gets into when it calls `coroutine.resume()` on some other coroutine: it is neither **running** nor **suspended**, since it can't be resumed when in this state.

running This is the state a coroutine is in when it's running

suspended Newly created coroutines are in this state, as well as coroutines that have suspended themselves (with `coroutine.yield()`)

dead The coroutine enters this state if the coroutine function returns, it is not possible to resume a dead coroutine

- Their status can be checked with `coroutine.status()`
- The real power comes from the `coroutine.yield()` function, which suspends the currently running coroutine and passes control back to the coroutine that caused it to run in the first place (with `coroutine.resume()`)
- `coroutine.resume()` runs in *protected mode*, so any error raised from within the coroutine will be returned by it, just like with `pcall()`
- The resume and yield functions can **exchange data**: an argument to any of them becomes a return value of the other
- *Symmetric coroutines* of other languages can be easily emulated in Lua
- Coroutines offer a great way to tackle the *producer-consumer problem* (the *who-has-the-main-loop* problem)
- They can also turn the caller/callee relationship inside out: now the callee can request from the caller (by resuming the caller)
- Coroutines offer a kind of *non-preemptive (cooperative)* multitasking, but since there is no parallelism involved, the code is easy to debug and there is no need for synchronization
- The cost of switching between coroutines is **really small** compared to switching between processes (as in UNIX pipes)
- They can be used to easily write iterators without having to worry about keeping a state
- The non-preemptive multitasking that they offer can be used to concurrently download files from the internet if *non-blocking sockets* are available

Complete Examples

- The *eight-queen puzzle* is solved with a configuration of eight queens on a chessboard in a way that no queen can attack another one
 - Any valid solution must have exactly one queen in each row
 - No two queens can be in the same column
 - No two queens can be directly diagonal to each other
- A *Markov-chain algorithm* can be used to generate pseudo-random text based on what words can follow a sequence of n previous words in a base text
- The resulting text is very, but not quite, random

Data Structures

- Tables in Lua are *the* data structure
- All other data structures can be implemented quite efficiently on top of Lua tables
- *Arrays* in Lua are implemented by simply indexing tables with integers
- They can grow as needed, elements that aren't used (equal to `nil`) don't take up space
- Indexing commonly starts at 1, but it can start at any other value
- *Matrices and multi-dimensional arrays* can be represented as either true *multi-dimensional arrays* or as *flat arrays*
- *Two-dimensional arrays* can be implemented with arrays by using a large array for the rows, containing one array per row which holds the actual data
- Allow for the most freedom (can represent *triangular matrices*)
- *Flat arrays* can also be used to represent matrices
- These can have an integer index that is composed of the two matrix indexes
- Otherwise, they can have a string index which is the concatenation of the two matrix indexes
- Either way, they can both easily represent matrices and are also efficient when working with *sparse matrices*: only matrix elements that are not `nil` take up memory
- Keys in tables have not intrinsic order, so iterations with `pairs()` happen in no particular order
- Linked lists are particularly easy to implement in Lua: they are simply a list with a reference to the next list
- Linked lists aren't used very often since they are not really necessary
- *Queues* and *double queues* can be implemented easily with tables that have a first and last index variables
- The first and last variables will continually increase when using the queues, but the available ~53 bits of integer precision (2^{53} representable integers) are unlikely to run out
- *Sets* and *bags* can be represented by storing the objects as the keys of a table (fast lookup)
- Objects in *sets* can either be in the set (then `set[obj]==true`) or not be in the set (then `set[name]==nil`)
- This behavior is achieved by using the object as indices of a table and setting its value to `true`
- *Bags* (also called **multisets**) are like **sets**, but there can be duplicates
- This is again trivial to implement with tables, where `bag[obj]` is either `nil` (then `obj` isn't in `bag`) or a number describing how many times `obj` is in `bag`
- **String buffers** are sometimes needed in Lua since strings are immutable
- To read a file chunk by chunk, all the chunks can first be stored in a table

and then finally put together with `table.concat()`, which optionally takes a string to use as separator

- File should, however, rather be read with `io.read("*a")`

Data Files and Persistence

- Writing data to file is easier than reading it back, since then it needs to be *parsed*
- Coding robust input routines is always difficult
- Lua started out as a data description language
- BibTeX was one of the inspirations for the constructor syntax in Lua
- We can use plain Lua to store data, which will look like this:

```
Person{
    name = "John Doe",
    age = 35,
    email = "john@example.com",
}
```

- `Person` both describes the data and represents a Lua function call, so we only have to define a sensible callback function and run the data file
- This is a *self-describing data format*, which means that it's easy to read and edit by hand
- Lua runs fast enough to store data like this, since data description has been one of the main applications of it
- To be able to write data which needs to be read back, it needs to be put in a known state, this process is called *serialization*
- We can do this recursively in Lua
- Floating-point numbers may lose precision when written and read back in decimal form, but we can use a hexadecimal format when writing: `string.format("%a", 0.4342)`
- Strings can also be properly escaped: `string.format("%q", str)`
- Tables can be serialized recursively **only** if they do not have *cycles* (where some parts of the table refer to other parts of the same table)
- To represent cycled tables, named tables are needed

Metatables and Metamethods

- *Metatables* allow us to change the behavior of a value when confronted with an undefined operation
- Whenever Lua tries to perform arithmetic operations on tables, it checks if any of them have a *metamethod* which defines this operation and runs it, otherwise it raises an error
- Tables and userdata have individual metatables, all other data types share one single metatable for all values of that type
- Lua always creates new tables without metatables
- We can use `setmetatable(t, mt)` to set or change the metatable of any type
- Any table can be the metatable of any value
- A group of related tables can share a metatable which describes their common behavior
- A table can be its own metatable so that it describes its own behavior
- Arithmetic metamethods are as follows:

```
__add addition (+)
__mul multiplication (*)
__sub subtraction (-)
__div division (/)
__unm negation (not)
__mod modulus (%)
__pow exponentiation (^)
__concat concatenation (..)
```

- When two tables are in an undefined expression, the left table's metamethod will be used if it exists, else the right table's metamethod will be used, and if that doesn't exist either, an error will be raised
- Metatables can also be used for *relational operators*:

```
__eq equality (==)
__lt less than comparison (<)
__le less than or equal (<=)
```
- Equality comparison doesn't work if objects have different types (always returns `false`)

- Some library functions also use metamethods to change their behaviour:

__tostring Used by `tostring()` to convert the table to a string

__metatable Returned by `getmetatable()` if it exists, and subsequent to setting this, `setmetatable()` will raise an error (as this marks the table as protected)

__pairs If defined, this is called by `pairs()` to get an iterator for the table (especially useful for proxy tables, where the actual data is not stored in the table)

__ipairs Just like **__pairs**, but for the `ipairs()` function

- The behaviour of tables can also be changed with metatables:

__index Called when a nonexistent index of the table is accessed. Can be either a metamethod (function) or a table, the latter is useful for using this to implement inheritance. Otherwise, it can also be used to change the default value of an empty table. This metamethod can be bypassed by using the `rawget()` function.

__newindex Called when a new table index is defined. This can be bypassed by using the `rawset()` function.

- The table access metamethods allow us to write a table that has no data of its own but simply proxies all access to another table, for example to track access or to block certain kinds of access (eg. *read-only*)

The Environment

- Lua keeps all global variables in a regular table, called the *global environment*
- This table can be manipulated like any other table
- It is globally accessible by the name `_G`
- Can be used to dynamically get the contents of variables, like so: `_G["varname"]`
- However, something such as `_G["io.read"]` won't work, since `io.read` is not a variable (only `io` is, and `read` is a member of it)
- Global variables do not need declarations
- We can change this behaviour if we like
- Since the global environment is a regular table, we can use metatables to change its behavior
- In a metamethod, `debug.getinfo(2, "S")` returns a table whose field `what` tells whether the function that called it is a main chunk, a regular Lua function or a C function
- In Lua, global variables do not need to be truly global (eg. they could be *nonlocal*)

- A *free name* is a name that is not bound to an explicit declaration, that is, it does not occur inside the scope of a local variable with that name
- The Lua compiler translates any free name `var` to `_ENV.var`
- This new `_ENV` variable is **not** a global variable
- Lua compiles any chunk in the presence of a predefined upvalue called `_ENV`, which is initialized by default with the global environment
- That is, Lua only has local variables, and any global variables are translated to local ones (`var` to `_ENV.var`, where `_ENV` is a local variable or upvalue)
- `_ENV` is, like `_G`, a plain regular (table) variable
- `_ENV` will refer to whatever `_ENV` variable is visible at that point in the code
- The assignment `_ENV=nil` will invalidate any direct access to global variables in the rest of the chunk
- The main use of `_ENV` is to change the environment used by a piece of code
- Again, we can change the behavior of `_ENV` with metatables
- The `load()` function has an optional fourth parameter that gives the value for `_ENV`
- Using that, we can limit or change the environment for external code
- We can also use `debug.setupvalue()` to change the upvalue for a compiled function
- When the function is the result of `load()` or `loadfile()`, Lua ensures that it has only one upvalue and that this upvalue is `_ENV`

Modules and Packages

- A *module* is some code (Lua or C) that can be loaded through `require()` and that creates and returns a table
- Everything that a module exports is defined inside this table, which works as a *namespace*
- Thus, we can manipulate modules like any other table
- Once a module is loaded, it will be reused by whatever part of the program requires it again
- The first step of the call `require "modname"` is to check in `package.loaded` whether the module is already loaded
- If not, it searches for a Lua file with the module name, which would be loaded with `loadfile()`
- If that doesn't work either, it looks for a C library with the module name
- If it finds a C library, it looks for the function `luaopen_modname()`
- Once it got a loader function, this is called with two arguments: the module name and the name of the file where it got the loader
- To find Lua files, Lua uses a couple of *templates* defined in `package.paths`
- The same applies the C libraries, but their paths are in `packages.cpaths`
- Writing a module is simple: create a table with all the functions that

should be exported in it and return it

- We can use some `_ENV` tricks to automatically build the table
- Lua also supports *submodules*, their names are formed with dots:
`module.submodule`
- To find submodules, Lua translates the name of the submodule into a path, like this: `/usr/local/lua/module/submodule.lua`
- For C libraries, `require()` looks for a function called `luaopen_module_submodule()` to load the submodule

Object-Oriented Programming

- A table in Lua is like an object in more than one sense
 - tables have a state
 - tables have an identity (a `self`) that is independent of their values
 - two objects (tables) with the same value are different objects
 - an object can have different values at different times
 - tables can, like objects, have their own operations
- To implement *methods* (functions that work on objects), a reference to the state of the object is needed
- Lua has special syntax for this:

```
Account = {balance=0} -- array object table
```

```
function Account:withdraw(amount)
    -- 'self' is the state of the object
    self.balance = self.balance - amount
end
```

- The new syntax uses the *colon operator* (`:`), which hides having to pass the state to methods away
- It adds an extra hidden parameter in a method definition and in a method call (eg. `Account:withdraw(20)` implicitly passes an argument `self` to the function)
- *Classes* in Lua can be implemented with tables and some metatables magic
- Because of the way metatables work in Lua, a `new()` method can create instances of a class and provide an interface to make subclasses:

```
function Account:new(o)
    o = o or {}
    setmetatable(o, self) -- setting metatable
```

```

        self.__index = self      -- allows subclassing
    return o
end

```

- With this kind of setup, any methods and default values will come from Account (due to the metatable), anything else will be stored in the table
- We can now override class methods in an object
- Also, like this any object acts as a class itself (you can call the `new()` method on any object of this class), this is how subclassing works:

```

-- standard account object
acc = Account:new{balance = 5}

-- subclass Account
SpecialAccount = Account:new()

-- special account object
sacc = SpecialAccount:new{balance = 10}

```

- Multiple inheritance needs some extra work to set up the metatable correctly
- The way multiple inheritance can be implemented is by defining a metamethod `__index` which looks for a requested method in all of the parents
- This means that multiple inheritance is slower than normal inheritance
- There are ways of restricting access to object data in Lua, but they are not used often
- One way is to use a proxy class which only allows access to methods but not private data

Weak Tables and Finalizers

- Lua does automatic memory management
- Lua automaticallz deletes objects that become garbage, using *garbage collection*
- Lua has no problems with cycles
- Sometimes the garbage collector needs your help
- *Weak tables* and *finalizers* are mechanisms that you can use in Lua to help the garbage collector

- *Weak tables* allow the collection of Lua objects that are still accessible to the program, while finalizers allow the collection of external objects that are not under control of the garbage collector
- A garbage collector can collect only what it can be sure is garbage
- Any object stored in a global variable is not garbage to Lua, even if it isn't used again
- Weak tables allow you to tell Lua that a reference should not prevent the reclamation of an object
- A *weak reference* is a reference to an object that is not considered by the garbage collector
- A *weak table* is a table whose entries are weak
- Tables have strong keys and values by default
- The weakness of a table is given by the field `__mode` of its metatable, which can be `k` for weak keys, `v` for weak values or a combination of those
- Only objects can be collected from a weak table
- Like numbers and booleans, strings are not collected from weak tables
- Weak tables can be used to *memoize* functions
- In Lua 5.2, a table with weak keys and string values is an *ephemeron table*, where the accessibility of a key controls the accessibility of the corresponding value
- A *finalizer* is a function associated with an object that is called when that object is about to be collected
- Finalizers are implemented through the metamethod `__gc`
- This metamethod needs to be present when setting the metatable to mark the object for finalization
- When the finalizer runs, it gets the object to be finalized as a parameter, this is called *transient resurrection*
- The finalizer could store the object in a global variable, *permanently resurrecting* it
- Finalizers are run in reverse order that the objects were marked for finalization in
- The finalizer for each object runs exactly once
- Objects with finalizers are collected in two phases:
 1. The collector detects that they are not reachable and calls their finalizers
 2. Next time the collector detects that the object is not reachable, it deletes the object

The Mathematical Library

- Comprises a standard set of mathematical functions
 - trigonometric functions (`sin`, `cos`, `tan`, `asin`, `acos`, ...)
 - exponentiation and logarithms (`exp`, `log`, `log10`)

- rounding functions (`floor`, `ceil`)
- `max` and `min`
- pseudorandom number generation functions (`random`, `randomseed`)
- Additionally, it contains some variables
 - `pi`, the mathematical constant pi
 - `huge`, the largest representable number
- All trigonometric functions use radians, but the functions `rad` and `deg` can be used to convert between them and degrees
- The `random` function returns a random real number between 0 and 1 when called without arguments
- When called with one argument, an integer `n`, it returns a random number between 1 and `n` inclusive
- The pseudorandom number generator needs to be seeded, this can be done with the current time with `math.randomseed(os.time())`

The Bitwise Library

- It is not easy to conciliate bitwise operations with floating-point numbers
- Lua 5.2 offers bitwise operations through a library: the `bit32` library
- The bitwise library operates on unsigned 32-bit integers
- It defines the following functions:

band(...) Binary *and* of the passed numbers

bor(...) Binary *or* of the passed numbers

bnot(n) Binary *not* (negation) of *n*

btest(...) Same as binary *and*, but returns `true` if the result is nonzero and `false` otherwise

lshift(n, a) Shift all bits of *n* to the left by a given amount *a*, filling empty spots with zero bits

rshift(n, a) Shift all bits of *n* to the right by a given amount *a*, filling empty spots with zero bits

arshift(n, a) Shift all bits of *n* to the right by a given amount *a* (to the left if that amount is negative), fill vacant bits on the left with copies of the last bit (the *signal bit*)

lrotate(n, a) Rotate the bits of *n* to the left by a given amount *a*

rrotate(n, a) Rotate the bits of *n* to the right by a given amount *a*

extract(n, p, w) Extract *w* bits (1 if no *w* parameter specified) of *n*, starting at bit *p* (bits are counted from the right, starting at 0)

replace(n, b, p, w) Replace w bits of n with the ones specified in b , starting at position p

- All functions will convert numbers to be in the range 0 to MAX, where max is $2^{32}-1$
- Numbers can be manually converted with the **bor()** or the **band()** function (by passing them as the sole argument)

The Table Library

- Comprises auxillary functions to manipulate tables as arrays
- Provides functions to insert and remove elements from lists, to sort the elements of an array and to concatenate all string in an array:

table.insert Inserts an element in a given position, moving up other elements to open space

table.remove Removes (and returns) an element from a given position in an array, moving down other elements to close space

table.sort Sorts an array, takes the array and optionally a sorting function, otherwise it uses the comparison operator ($<$) for sorting. This function can't sort tables, only sequences (arrays)!

table.concat Takes a list of strings and returns the result of concatenating all these strings. It can take an optional separator string to separate the strings from the array

The String Library

- All functions can be found in the **string** table as well as in the metatable of all strings
- The basic string functions are:

string.len Returns the length of a string, equivalent to the length operator ($\#$)

string.rep Repeats a string a given amount of times

string.lower Uses the current locale setting to turn all characters of a given string into lower-case characters

string.upper Uses the current locale setting to turn all characters of a given string into upper-case characters

string.sub From a given string, return a substring of that starting and ending at given indexes

string.char Converts a number to whatever character it stands for in the current locale

string.byte Converts a string into it's internal numeric representation

string.format Format a string using a *format string* similar to the ones the C function **printf** accepts

- The string library offers additional functions for pattern matching, string searching and substitution
- Patterns in Lua are not full regular expressions but they are very similar to them with a few differences
- These are some of the constructs Lua supports:
 - *character classes*: **%a** (letters), **%c** (control characters), **%d** (digits), **%l** (lower-case letters), **%u** (upper-case letters), **%p** (punctuation letters), **%s** (space characters), **%w** (alpha-numeric characters) and **%x** (hexadecimal digits)
 - *magic characters*: pretty much just like in regular expressions, however **%** is used as the escape character (instead of the backslash)
 - *character groups* are done by enclosing the list of allowed characters in square brackets (this will match any English vowel: **[aeiou]**)
 - *groups* (called *captures* here) are enclosed in round brackets, like so: **(%u%l+)**. Pattern matching functions will return all captures individually, and as opposed to regular expressions, there can't be a variable amount of them (eg. they can't be followed by a star or a plus)
 - *lazy repetitions* can be denoted with a minus, like so: **%u[%a%s]-%**, the minus means the shortest possible sequence will be matched
- Lua pattern matching facilities are not meant to work with UTF-8 encoded text, while they can be made to work some things will not work as intended, for example the character classes won't work for all characters
- Unicode text can be represented in Lua strings, but the functions that the string library offers don't all work on them: **string.reverse**, **string.byte**, **string.char**, **string.upper** and **string.lower** all shouldn't be used on UTF-8 strings

The I/O Library

- Lua offers two different models for file manipulation
- The *simple I/O model* does all it's operations on two current files: the standard input (*stdin*) and the standard output (*stdout*)

- By default, *stdin* and *stdout* are connected to the console from which the program is executed
- Operations such as `io.write()` operate on the standard input and output files
- These current files can be changed with the `io.input()` and `io.output()` functions
- There are two main functions to work with the simple model:
 - `io.write` gets an arbitrary number of string arguments and writes them to the current output file, similar to the `print` function, but offers more control over the output.
 - `io.read` reads strings from the current input file, and takes arguments describing what to read:
 - `*a` to read the whole file
 - `*l` to read a single line (without newline)
 - `*L` to read a single line (with newline)
 - `*n` to read a number, skipping any spaces prior to it or `nil` when it can't find one
 - `num` reads a string with up to *num* characters
- The *complete I/O model* offers more control and multiple open files with something called *file handles*, which are equivalent to `FILE*` streams in C
- Every open file has a *file handle* associated with it, which is an object that has methods for manipulating the file
- The `io.open()` function can be used to open a file and get a handle for it, it takes as arguments the name of the file and in which *mode* to open it: possible are `r` for reading, `w` for writing and `a` for appending and an optional `b` to open binary files
- Just like other Lua library functions, it returns `nil` plus an error message and an error number in case of an error
- After having opened a file `f`, there are read methods (`f:read()`) and write methods (`f:write()`) available with the same semantics as those from the simple I/O model
- There are predefined file handles: `io.stdin`, `io.stdout` and `io.stderr` for the three standard C streams
- The complete model and the simple model can be mixed, for example `io.input():write()` does the same as `io.write()`
- To read big files in Lua it is advisable to use a relatively large buffer size of efficiency

- On Windows systems care must be taken to open files in the correct mode: binary files must be opened in *binary mode* to avoid their contents being changed while reading
- Binary data is usually read with the `*a` flag or with a given size of n bytes
- `io.tmpfile()` returns a handle for a temporary file, open in read/write mode
- For a file handle `f`, `f:flush()` executes all pending writes to the file
- The `setvbuf` method sets the buffering mode of a stream, it can be set to "no" (no buffering), "full" (the file is only written when the buffer is full or it's flushed) or "line" (it is buffered until a newline is output).
- The `seek()` method can be set with the `set(whence, offset)` method, the `whence` parameter is a string that specifies how to interpret the offset: "set" interprets strings from the beginning of the file, "cur" interprets them from the current position, and "end" interprets them relative to the end
- Calling `seek()` on a file pointer (like `f:seek()`) returns the current position in the file

The Operating System Library

- It is defined in the table `os`
- Includes functions for file manipulation, getting the current date and time, and other facilities related to the operating system
- Because Lua is written in ANSI C, it uses only the functionality that the ANSI standard offers
- For date and time, the `os` library offers two functions
- The `os.time()` function returns the current date and time, coded as a number (usually a UNIX timestamp)
- When called with a table, it returns the number representing the date and time described by the table
- Such a *date table* has the following fields: `year`, `month`, `day`, `hour`, `min`, `sec`, `isdst` (true if daylight saving is on)
- Note that the result of the `os.time()` call depends on the system's time zone
- The `os.date()` function is a kind of a reverse of the `os.time()` function: it converts the number describing the date and time into some higher-level representation

- It's first parameter is a format string describing the representation we want
- To produce a table, we use the format string `*t`, like so: `os.date("*t", 906000490)`
- Format strings have a similar syntax to that of `string.format()`, however there are more possible values
- For timing, the `os` library offers `os.clock()`, which returns the number of seconds of CPU time have passed for the program
- Other systems calls that the library supports are leaned on the called available in standard C:

`os.exit()` terminates the execution of the program, taking as parameter the return status of the program, and as second argument whether to close the Lua state (and calling finalizers) or not

`os.getenv()` gets the value of an environment variable, taking as argument the name of the variable and returning it's value

`os.execute()` runs a system command

`os.setlocale()` sets the current locale used by the Lua program

- Care must be taken to set the locale of the Lua program to the standard C locale when creating pieces of code from within Lua

The Debug Library

- The debug library offers all the primitives needed to build a debugger for Lua
- Should be used sparingly because it lacks performance and it breaks truths about the language
- It offers two kinds of functions:
 - *introspective functions*, which allow us to inspect several aspects of the running program
 - *hooks*, which allow us to trace the execution of a program
- An important concept in the debug library is the *stack level*, which refers to a particular function in the calling hierarchy

`debug.getinfo` main introspective function in the debug library, it offers insight into a given function but is quite slow

`debug.traceback` constructs a traceback (a graph of the calling hierarchy) and returns it as a string

`debug.getlocal` allows us to inspect the local variables of any active function as well as the parameters passed to it

`debug.setlocal` like `getlocal`, but modifies a local variables instead of just returning it

`debug.getupvalue` grants access to the nonlocal variables of a function and does not need the function to be active

`debug.setupvalue` modifies a closure's upvalues (nonlocal variables)

- Introspection functions can also optionally take a coroutine as parameter to allow inspection of their states
- There are four kinds of hooks in the debug library:
 - *call* events when Lua calls a function
 - *return* events when a function returns
 - *line* events when Lua starts executing a new line of code
 - *count* events after a given number of instructions
- `debug.sethook` registers a hook with Lua so it will be called when the requested event takes place
- `debug.debug` starts an interactive debugging console
- The hook mechanism can be used to profile code

An Overview of the C API

About the C API

- Lua is an *embedded language*
- There are two kinds of interactions between Lua and C code:
 - As a **stand-alone language** that can be extended with native C code (the C code is the library), this makes Lua an *extensible language*
 - As a **library** to integrate into C projects to extend them (Lua is the library), this makes Lua an *extension language*
- The difference between the two kinds of interactions is the language which has the control
- The Lua interpreter uses Lua as a library
- The C API is the set of functions that allow C code to interact with Lua, it comprises a set of functions to:
 - read and write Lua global variables
 - call Lua functions
 - run pieces of Lua code
 - register C functions to be callable from within Lua

- The C API follows the *modus operandi* of C, which means that we must care about several inconveniences, including:
 - type checking
 - error recovery
 - memory-allocation errors
- Anything that can be done within Lua can also be done with the C API, however it may be more lengthy (common tasks may involve several API calls)
- The major component in the communication between Lua and C is an omnipresent virtual *stack*

The Header Files

- The file `lua.h` defines the basic functions provided by Lua, it includes functions to do the following:
 - create a new Lua environment
 - invoke Lua functions (eg. `lua_pcall`)
 - read and write global variables in the Lua environment
- All functions in `lua.h` are prefixed with `lua_`
- The header file `luauxlib.h` defines the functions provided by the auxiliary library (*auxlib*)
- All functions in `luauxlib.h` are prefixed with `luaL_`
- They use the basic functions defined in `lua.h` to provide a higher abstraction level
- The Lua library keeps all its state in the dynamic structure `lua_State`
- New states can be created by the `luaL_newstate` function
- The standard libraries are not loaded by default, this can be done with `luaL_openlibs`
- The function `luaL_loadstring` compiles Lua code, and pushes the resulting function to the stack, returning nonzero on error.

Extending Your Application

- An important use of Lua is as a *configuration language*
- You can use Lua programs as very flexible configuration files
- You must use the Lua API to parse the file and then get the values of the global setting variables
- Using `lua_pcall` to run the files means that errors are easily caught
- With `lua_getglobal`, global variables can be accessed
- Care must be taken to ensure type safety when reading Lua variables from C

- Such configuration scripts allow complex configurations, for example varying values depending on environment variables
- This results in more flexible programs
- Complex objects can also be read using `lua_gettable` (which takes the key to read from the table from the stack) or `lua_getfield` (which takes the key as parameter)
- For setting a table's fields, the functions `lua_settable` and `lua_setfield` are available
- New tables can be created on the stack with `lua_newtable`

Calling a Lua function from C

- To call a Lua function from within C, the following steps must be done:
 1. push the function to the stack (for example with `lua_getglobal` if the function is globally visible)
 2. push the arguments to the call
 3. use `lua_pcall` to do the actual call
- `lua_pcall` takes four arguments: the Lua state, the number of arguments passed, the number of return values expected and optionally the stack index of an error-handling function
- `lua_pcall` leaves the requested amount of return values on the stack, either trimming them if the function returns too many or adding extra nils
- When multiple results are returned, the first result is pushed first and the last one last, leaving it at the top of the stack
- If there is any error while `lua_pcall` is running, it returns one of these error codes to indicate so: `LUA_OK` : no error occurred

`LUA_ERRRUN` a normal error occurred

`LUA_ERRMEM` a memory allocation error occurred

`LUA_ERRERR` an error occurred while running the error message handler

`LUA_ERRGCMM` an error occurred while running a garbage-collector metamethod

- Additionally, when an error occurred there will be no function return values on the stack, instead there will be an error message describing what went wrong
- If there was a message handler set when running `lua_pcall`, it will be run to create the error message, else the standard message is used